

CAS CS 460/660

Introduction to Database Systems

Query Evaluation I

Slides from UC Berkeley

Introduction

- We've covered the basic underlying storage, buffering, and indexing technology.

- Now we can move on to query processing.

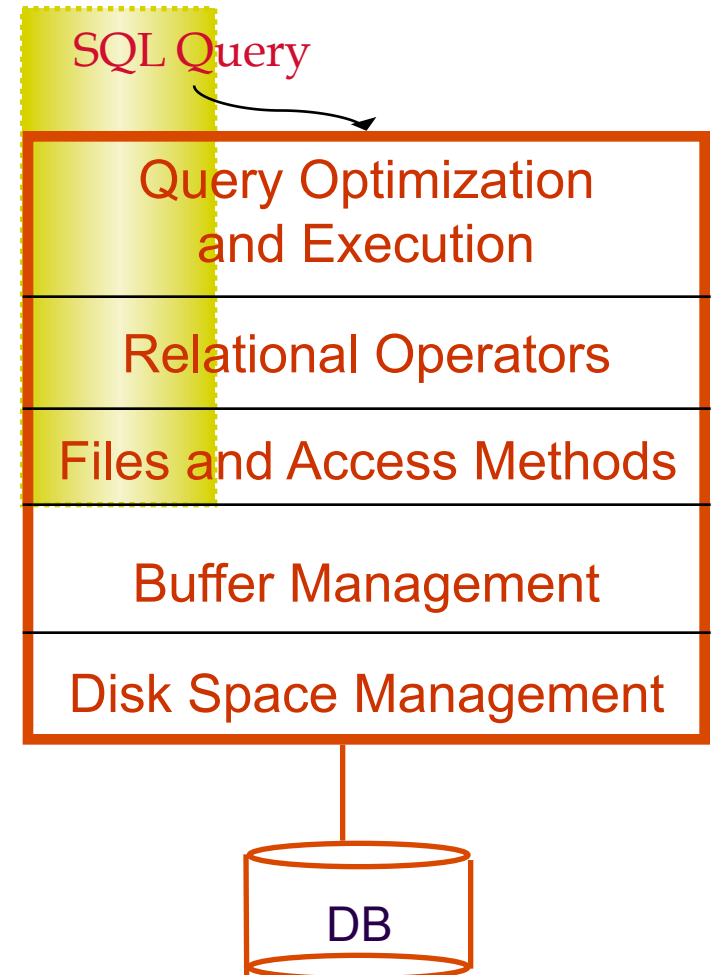
- Some database operations are **EXPENSIVE**

- Can greatly improve performance by being “smart”

- e.g., can speed up 1,000x over naïve approach

- Main weapons are:

1. clever implementation techniques for operators
2. exploiting “equivalencies” of relational operators
3. using statistics and cost models to choose among these.



Cost-based Query Sub-System

Queries

```
Select *  
From Blah B  
Where B.blah = blah
```

Query Parser

Query Optimizer

Plan
Generator

Plan Cost
Estimator


Catalog Manager

Schema

Statistics

Query Plan Evaluator


Usually there is a
heuristics-based
rewriting step before
the cost-based steps.



Query Processing Overview


- The *query optimizer* translates SQL to a special internal “language”
 - Query Plans
- The *query executor* is an *interpreter* for query plans
- Think of query plans as “box-and-arrow” *dataflow* diagrams
 - Each box implements a *relational operator*
 - Edges represent a flow of tuples (columns as specified)
 - For single-table queries, these diagrams are straight-line graphs

SELECT DISTINCT name, gpa
FROM Students



Query Optimization

- A deep subject, focuses on multi-table queries
 - ↗ We will only need a cookbook version for now.
- Build the dataflow bottom up:
 - ↗ Choose an Access Method (HeapScan or IndexScan)
 - Non-trivial, we'll learn about this later!
 - ↗ Next apply any WHERE clause filters
 - ↗ Next apply GROUP BY and aggregation
 - Can choose between sorting and hashing!
 - ↗ Next apply any HAVING clause filters
 - ↗ Next Sort to help with ORDER BY and DISTINCT
 - In absence of ORDER BY, can do DISTINCT via hashing!



Iterators

- The relational operators are all subclasses of the class `iterator`:

```
class iterator {  
    void init();  
    tuple next();  
    void close();  
    iterator inputs[];  
    // additional state goes here  
}
```



iterator

- Note:
 - Edges in the graph are specified by inputs (max 2, usually 1)
 - Encapsulation: any iterator can be input to any other!
 - When subclassing, different iterators will keep different kinds of state information

Example: Scan

```
class Scan extends iterator {  
    void init();  
    tuple next();  
    void close();  
    iterator inputs[1];  
    bool_expr filter_expr;  
    proj_attr_list proj_list;  
}
```

■ init():

- Set up internal state
- call init() on child – often a file open

■ next():

- call next() on child until qualifying tuple found or EOF
- keep only those fields in “proj_list”
- return tuple (or EOF -- “End of File” -- if no tuples remain)

■ close():

- call close() on child
- clean up internal state

Note: Scan also applies “selection” filters and “projections”
(without duplicate elimination)

Example: Sort

```
class Sort extends iterator {  
    void init();  
    tuple next();  
    void close();  
    iterator inputs[1];  
    int numberOfRuns;  
    DiskBlock runs[];  
    RID nextRID[];  
}
```

■ init():

- ✚ generate the sorted runs on disk
- ✚ Allocate runs[] array and fill in with disk pointers.
- ✚ Initialize numberOfRuns
- ✚ Allocate nextRID array and initialize to NULLs

■ next():


- ✚ nextRID array tells us where we're "up to" in each run
- ✚ find the next tuple to return based on nextRID array
- ✚ advance the corresponding nextRID entry
- ✚ return tuple (or EOF -- "End of File" -- if no tuples remain)

■ close():

- ✚ deallocate the runs and nextRID arrays

Streaming through RAM

- Simple case: “Map”. (assume many **records** per disk **page**)
 - ↗ Goal: Compute $f(x)$ for each record, write out the result
 - ↗ Challenge: minimize RAM, call read/write rarely
- Approach
 - ↗ Read a chunk from INPUT to an *Input Buffer*
 - ↗ Write $f(x)$ for each item to an *Output Buffer*
 - ↗ When Input Buffer is consumed, read another chunk
 - ↗ When Output Buffer fills, write it to OUTPUT
- Reads and Writes are **not coordinated** (i.e., not in lockstep)
 - ↗ E.g., if $f()$ is Compress(), you read many chunks per write.
 - ↗ E.g., if $f()$ is DeCompress(), you write many chunks per read.




Rendezvous

- Streaming: one chunk at a time. Easy.
- But some algorithms need certain items to be co-resident in memory
 - ↗ not guaranteed to appear in the same input chunk
- *Time-space Rendezvous*
 - ↗ in the same place (RAM) at the same time
- There may be many combos of such items

Divide and Conquer


- *Out-of-core* algorithms orchestrate rendezvous.
- Typical RAM Allocation:
 - ↗ Assume B pages worth of RAM available
 - ↗ Use 1 page of RAM to read into
 - ↗ Use 1 page of RAM to write into
 - ↗ $B-2$ pages of RAM as workspace



Divide and Conquer

■ Phase 1


- ↗ “streamwise” *divide* into $N/(B-2)$ megachunks
- ↗ output (write) to disk one megachunk at a time



Divide and Conquer


■ Phase 2

- ↗ Now megachunks will be the input
- ↗ process each *megachunk* individually.



Sorting: 2-Way

- Pass 0:
 - read a page, sort it, write it.
 - only one buffer page is used
 - a repeated “batch job”



Sorting: 2-Way (cont.)

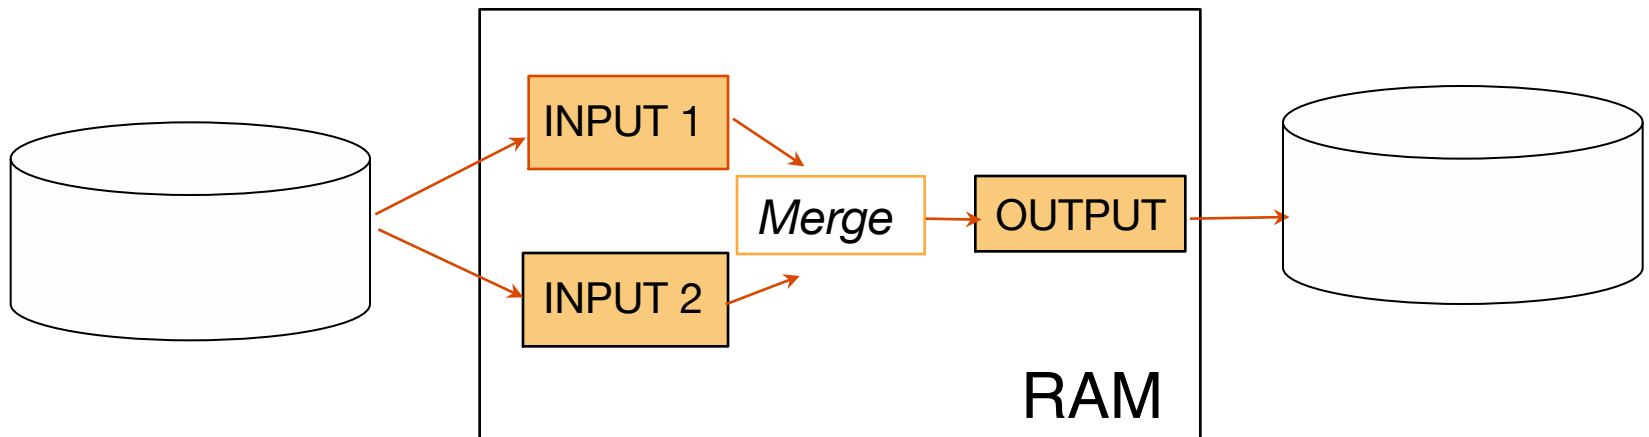
■ Pass 1, 2, 3, ..., etc. (merge):

↗ requires **3 buffer pages**

- note: this has nothing to do with double buffering!

↗ *merge pairs* of runs into runs twice as long

↗ a streaming algorithm, as in the previous slide!



Two-Way External Merge Sort

- **Sort subfiles and Merge**

- How many passes?

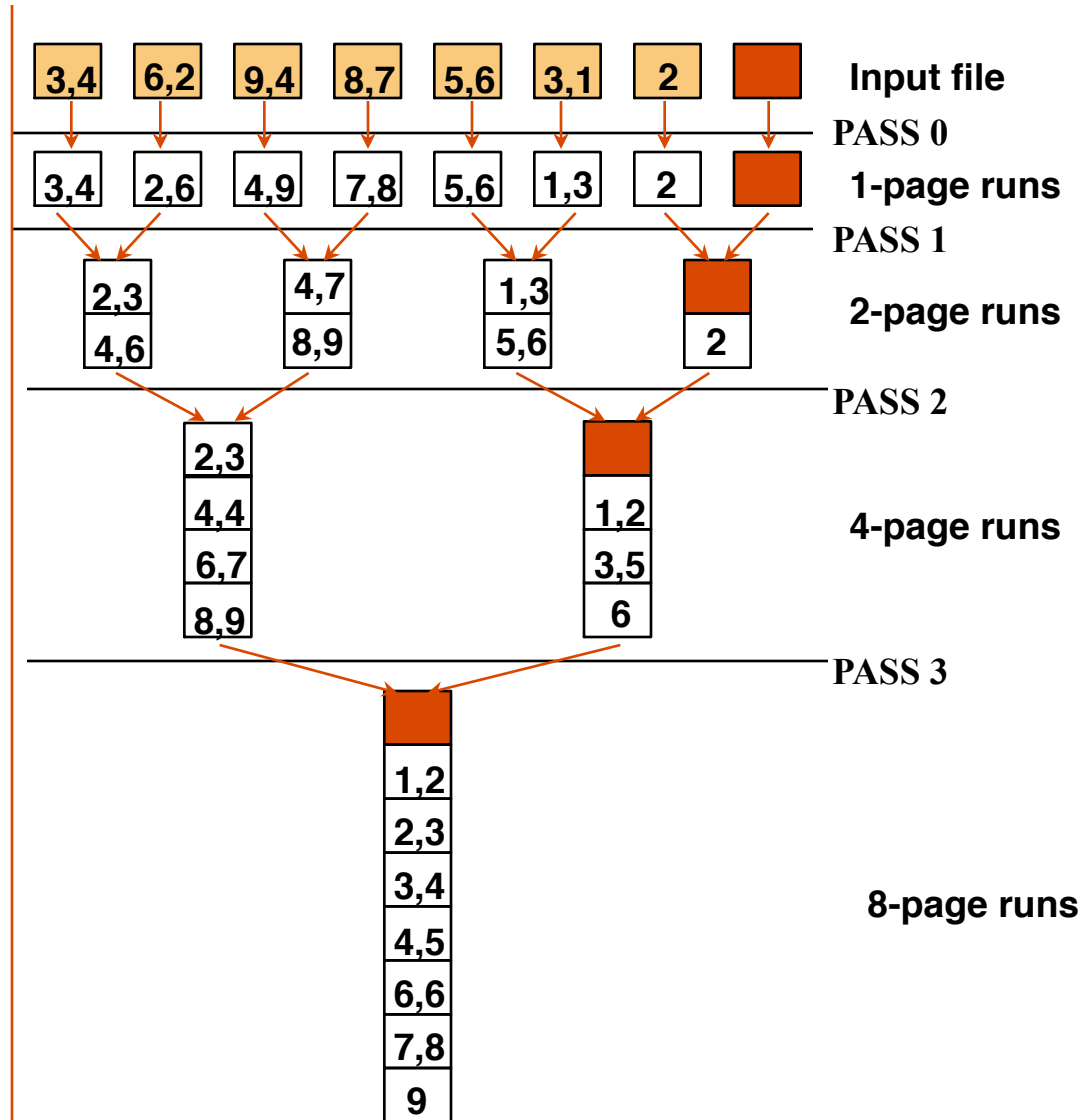
- N pages in the file
=> the number of passes =

$$\lceil \log_2 N \rceil + 1$$

- Total I/O cost? (reads + writes)

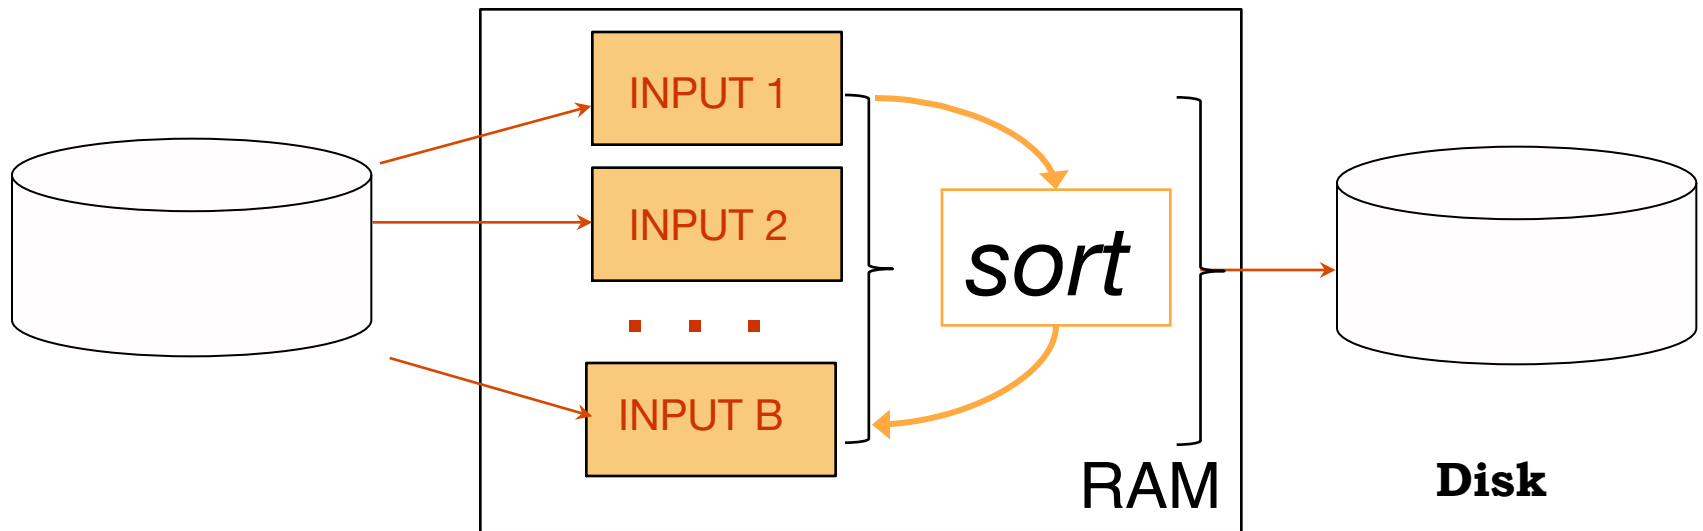
- Each pass we read + write each page in file. So total cost is:

$$2N(\lceil \log_2 N \rceil + 1)$$



General External Merge Sort

- More than 3 buffer pages. How can we utilize them?
- To sort a file with N pages using B buffer pages:
 - ↗ Pass 0: use B buffer pages. Produce $\lceil N / B \rceil$ sorted runs of B pages each.

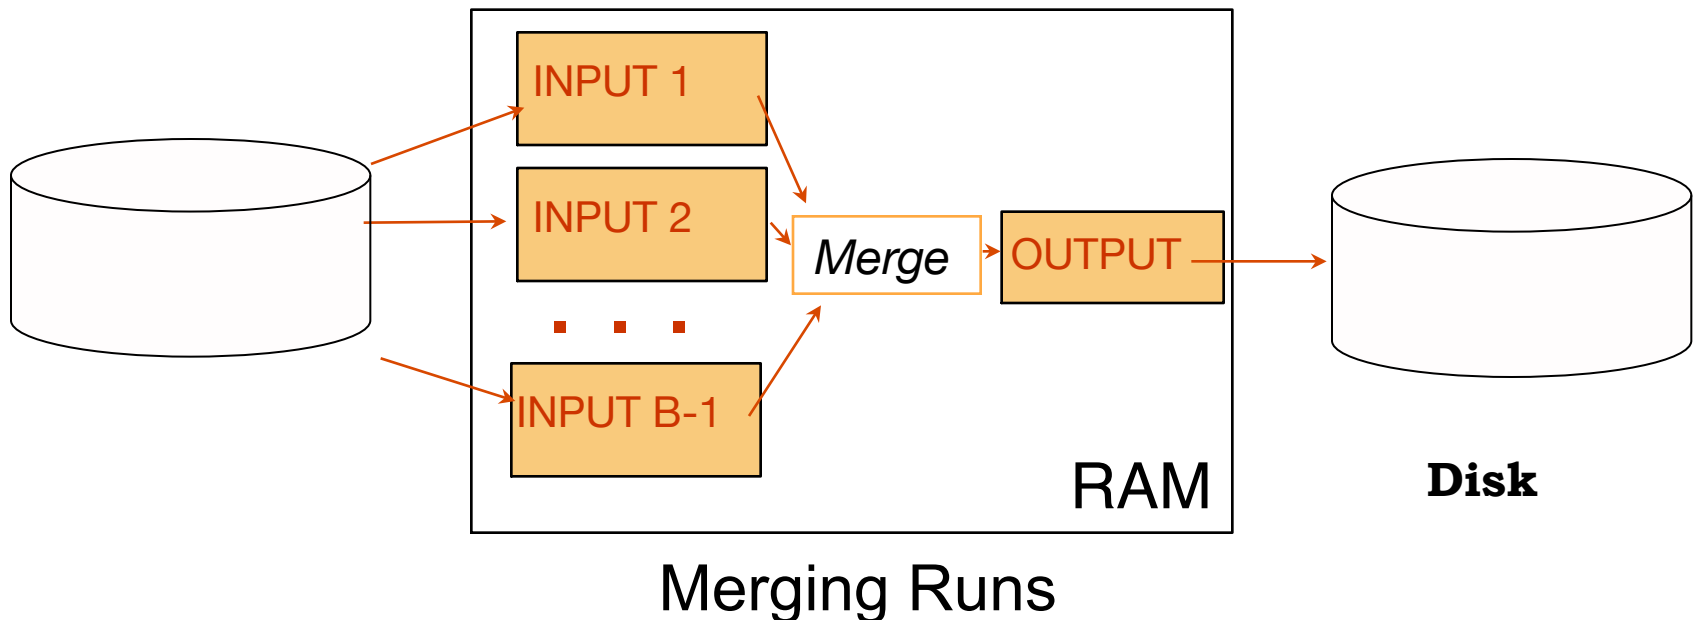


Pass 0 – Create Sorted Runs

General External Merge Sort

Pass 1, 2, ..., etc.: merge $B-1$ runs.

Creates runs of $(B-1) \times$ size of runs from previous pass.



Cost of External Merge Sort

- Number of passes: $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- Cost = $2N * (\text{\# of passes})$
- E.g., with 5 buffer pages, to sort 108 page file:
 - ↗ Pass 0: $\lceil 108 / 5 \rceil = 22$ sorted runs of 5 pages each (last run is only 3 pages)
 - ↗ Pass 1: $\lceil 22 / 4 \rceil = 6$ sorted runs of 20 pages each (last run is only 8 pages)
 - ↗ Pass 2: 2 sorted runs, 80 pages and 28 pages
 - ↗ Pass 3: Sorted file of 108 pages

Formula check: $1 + \lceil \log_4 22 \rceil = 1 + 3 \rightarrow \underline{4 \text{ passes}} \quad \checkmark$

of Passes of External Sort

(I/O cost is $2N$ times number of passes)

N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

Memory Requirement for External Sorting

- How big of a table can we sort in two passes?
 - ↗ Each “sorted run” after Phase 0 is of size B
 - ↗ Can merge up to B-1 sorted runs in Phase 1
- Answer: $B(B-1)$.
 - ↗ Sort N pages of data in about \sqrt{N} space

Alternative: Hashing

■ Idea:

- ↗ Many times we don't require order
- ↗ E.g.: removing duplicates
- ↗ E.g.: forming groups

■ Often just need to *rendezvous* matches

■ Hashing does this

- ↗ And may be cheaper than sorting! (Hmmm...!)
- ↗ But how to do it out-of-core??

Divide

- Streaming Partition (divide):
Use a hash f'n h_p to stream records to disk partitions
 - ↗ All matches rendezvous in the same partition.
 - ↗ *Streaming* alg to create partitions on disk:
 - “Spill” partitions to disk via output buffers

Divide & Conquer

■ Streaming Partition (divide):

Use a hash function h_p to stream records to **disk-based** partitions

✚ All matches rendezvous in the same partition.

✚ *Streaming* alg to create partitions on disk:

- “Spill” partitions to disk via output buffers

■ ReHash (conquer):

Read partitions into **RAM-based** hash table one at a time, using hash function h_r


✚ Then go through each bucket of this hash table to achieve rendezvous in RAM

■ Note: Two different hash functions

✚ h_p is coarser-grained than h_r


Two Phases

■ Partition:




Two Phases


■ Partition:



■ Rehash:



Cost of External Hashing



$$\text{cost} = 4 * N \text{ IO's}$$

Memory Requirement

■ How big of a table can we hash in two passes?

- ↗ B-1 “partitions” result from Phase 0

- ↗ Each should be no more than B pages in size

- ↗ Answer: $B(B-1)$.

 - We can hash a table of size N pages in about \sqrt{N} space

- ↗ Note: assumes hash function distributes records evenly!

■ Have a bigger table? Recursive partitioning!

- ↗ How many times?

 - Until every partition fits in memory !! ($\leq B$)

**How does this compare with
external sorting?**

So which is better ??

■ Simplest analysis:

- ↗ Same memory requirement for 2 passes
- ↗ Same I/O cost
- ↗ But we can dig a bit deeper...

■ Sorting pros:

- ↗ Great if input already sorted (or almost sorted) w/heapsort
- ↗ Great if need output to be sorted anyway
- ↗ Not sensitive to “data skew” or “bad” hash functions

■ Hashing pros:

- ↗ For duplicate elimination, scales with # of values
 - Not # of items! We'll see this again.
- ↗ Can exploit extra memory to reduce # IOs (stay tuned...)

Summing Up 1

- Unordered collection model
- Read in chunks to avoid fixed I/O costs
- Patterns for Big Data
 - ↗ Streaming
 - ↗ Divide & Conquer
 - ↗ also Parallelism (but we didn't cover this here)

Summary Part 2

- Sort/Hash Duality
 - ↗ Sorting is Conquer & Merge
 - ↗ Hashing is Divide & Conquer
- Sorting is overkill for rendezvous
 - ↗ But sometimes a win anyhow
- Sorting sensitive to internal sort alg
 - ↗ Quicksort vs. HeapSort
 - ↗ In practice, QuickSort tends to be used
- Don't forget double buffering (with threads)